

Brake & Turn Companion 01 — Living Checklist

Version 1.0 · Public Domain (CC0)

A. System Purpose & Fit

- Is the system truly necessary? Does it reduce harm or accelerate good?
- Is there a safer, simpler, or non-AI approach that meets the need?
- Are goals clearly defined, measurable, and time-bounded?

B. Data & Dignity

- Do we have a lawful basis, consent, or a legitimate interest for each dataset?
- Is sensitive data minimized, encrypted, and access-controlled?
- Are people able to opt-out, correct data, and see how it's used?

C. Energy & Footprint

- Is training/inference energy measured and published?
- Is compute scheduled on lower-carbon energy and right-sized for need?
- Are lifecycle impacts (cooling, e-waste, water) considered and mitigated?

D. Model Safety & Governance

- Are model risks (bias, hallucinations, leakage, jailbreaks) assessed and mitigated?
- Are outputs rate-limited, logged, and reviewed where harm is plausible?
- Is there a rollback plan if harms or regressions are detected?

E. Human Oversight & Labor

- Is human review required for sensitive, financial, legal, or medical decisions?
- Are end-users trained on limitations, escalation paths, and safe usage?
- Does the system respect labor (annotation, moderation) with fair pay?

F. Culture & Continuity

- Does the system regenerate value for humans, nature, and culture?
- Are outcomes evaluated not just for profit but for life-affirming impact?
- Is learning continuous, with feedback loops embedded?

Review this checklist quarterly. Update your assessments in alignment with community and environmental shifts.

When in doubt, ask: “Does this system serve life, or consume it?”



Scan resources → <https://cc2397.github.io/brakeandturn/companions/living-checklist/>